

Stock_Price_Prediction_Using_LSTM

Import Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: df = pd.read_csv("Stocks_dataset.csv")
df.head()
```

Out[2]:

	symbol	date	close	high	low	open	volume	adjClose	adjHigh	adjLow	adjOpen	adjVolume	di
0	GOOG	2016-06-14 00:00:00+00:00	718.27	722.47	713.1200	716.48	1306065	718.27	722.47	713.1200	716.48	1306065	
1	GOOG	2016-06-15 00:00:00+00:00	718.92	722.98	717.3100	719.00	1214517	718.92	722.98	717.3100	719.00	1214517	
2	GOOG	2016-06-16 00:00:00+00:00	710.36	716.65	703.2600	714.91	1982471	710.36	716.65	703.2600	714.91	1982471	
3	GOOG	2016-06-17 00:00:00+00:00	691.72	708.82	688.4515	708.65	3402357	691.72	708.82	688.4515	708.65	3402357	
4	GOOG	2016-06-20 00:00:00+00:00	693.71	702.48	693.4100	698.77	2082538	693.71	702.48	693.4100	698.77	2082538	

```
In [3]: #Shape of Data
df.shape
```

Out[3]: (1258, 14)

```
In [4]: #Statistical Descriptipion
df.describe()
```

Out[4]:

	close	high	low	open	volume	adjClose	adjHigh	adjLow	adjO
count	1258.000000	1258.000000	1258.000000	1258.000000	1.258000e+03	1258.000000	1258.000000	1258.000000	1258.000
mean	1216.317067	1227.430934	1204.176430	1215.260779	1.601590e+06	1216.317067	1227.430936	1204.176436	1215.260
std	383.333358	387.570872	378.777094	382.446995	6.960172e+05	383.333358	387.570873	378.777099	382.446
min	668.260000	672.300000	663.284000	671.000000	3.467530e+05	668.260000	672.300000	663.284000	671.000
25%	960.802500	968.757500	952.182500	959.005000	1.173522e+06	960.802500	968.757500	952.182500	959.005
50%	1132.460000	1143.935000	1117.915000	1131.150000	1.412588e+06	1132.460000	1143.935000	1117.915000	1131.150
75%	1360.595000	1374.345000	1348.557500	1361.075000	1.812156e+06	1360.595000	1374.345000	1348.557500	1361.075
max	2521.600000	2526.990000	2498.290000	2524.920000	6.207027e+06	2521.600000	2526.990000	2498.290000	2524.920

```
In [5]: #Data Summary
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1258 entries, 0 to 1257
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   symbol          1258 non-null   object
1   date            1258 non-null   object
2   close           1258 non-null   float64
3   high            1258 non-null   float64
4   low             1258 non-null   float64
5   open            1258 non-null   float64
6   volume          1258 non-null   int64
7   adjClose        1258 non-null   float64
8   adjHigh         1258 non-null   float64
9   adjLow          1258 non-null   float64
10  adjOpen         1258 non-null   float64
11  adjVolume       1258 non-null   int64
12  divCash         1258 non-null   float64
13  splitFactor     1258 non-null   float64
dtypes: float64(10), int64(2), object(2)
memory usage: 137.7+ KB
```

```
In [6]: #Check Null values
df.isnull().sum()
```

```
Out[6]: symbol      0
date              0
close             0
high             0
low              0
open             0
volume           0
adjClose         0
adjHigh          0
adjLow           0
adjOpen          0
adjVolume        0
divCash          0
splitFactor      0
dtype: int64
```



```
In [7]: #No null values in data
df = df[['date', 'open', 'close']] #Require columns extracting
df['date'] = pd.to_datetime(df['date'].apply(lambda x: x.split()[0])) #Convert to datetime dtype
df.set_index('date', drop=True, inplace=True)
df.head()
```

```
Out[7]:
```

	open	close
date		
2016-06-14	716.48	718.27
2016-06-15	719.00	718.92
2016-06-16	714.91	710.36
2016-06-17	708.65	691.72
2016-06-20	698.77	693.71

```
In [11]: # Now we plotting open and closing price on date index
fig, ax = plt.subplots(1,2, figsize=(20,7))
ax[0].plot(df['open'], label = 'open', color = 'green')
ax[0].set_xlabel('Data', size=13)
ax[0].set_ylabel('Price', size=13)
ax[0].legend()
plt.show()

ax[1].plot(df['close'], label = 'Close', color = 'red')
ax[1].set_xlabel('Data', size=13)
ax[1].set_ylabel('Price', size=13)
ax[1].legend()
plt.show()
```



Data Preprocessing

```
In [12]: # we'll normalizing all the values of all columns using MinMaxScaler
from sklearn.preprocessing import MinMaxScaler
mms = MinMaxScaler()
df[df.columns] = mms.fit_transform(df)
df.head()
```

Out[12]:

	open	close
date		
2016-06-14	0.024532	0.026984
2016-06-15	0.025891	0.027334
2016-06-16	0.023685	0.022716
2016-06-17	0.020308	0.012658
2016-06-20	0.014979	0.013732

```
In [13]: #split data into Training and testing
training_size = round(len(df) * 0.75) #75% for training
training_size
```

Out[13]: 944

```
In [14]: training_data = df[:training_size]
testing_data = df[training_size:]
training_data.shape, testing_data.shape
```

Out[14]: ((944, 2), (314, 2))

```
In [15]: # we'll create sequence of data for training and testing

def create_sequence(dataset):
    sequence = []
    labels = []

    start_idx = 0

    for stop_idx in range(50, len(dataset)): #selecting 50 rows at a time
        sequence.append(dataset.iloc[start_idx:stop_idx])
        labels.append(dataset.iloc[stop_idx])
        start_idx += 1
    return (np.array(sequence), np.array(labels))
```

```
In [16]: train_seq, train_label, = create_sequence(training_data)
test_seq, test_label = create_sequence(testing_data)
train_seq.shape, train_label.shape, test_seq.shape, test_label.shape
```

```
Out[16]: ((894, 50, 2), (894, 2), (264, 50, 2), (264, 2))
```

Create LSTM Model

```
In [17]: from keras.models import Sequential
from keras.layers import Dense, Dropout, LSTM, Bidirectional
```

```
In [18]: #import Sequential from keras.models
model = Sequential()
#import Dense, Dropout, LSTM, Bidirectional from keras.layers
model.add(LSTM(units=50, return_sequences=True, input_shape=(train_seq.shape[1], train_seq.shape[2])

model.add(Dropout(0.1))
model.add(LSTM(units=50))

model.add(Dense(2))

model.compile(loss = 'mean_squared_error', optimizer = 'adam', metrics = ['mean_absolute_error'])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
lstm (LSTM)	(None, 50, 50)	10600
dropout (Dropout)	(None, 50, 50)	0
lstm_1 (LSTM)	(None, 50)	20200
dense (Dense)	(None, 2)	102
=====		
Total params: 30902 (120.71 KB)		
Trainable params: 30902 (120.71 KB)		
Non-trainable params: 0 (0.00 Byte)		

```
In [20]: # fitting the model by iterating the dataset over 100 times(100 epochs)
model.fit(train_seq, train_label, epochs = 100, validation_data = (test_seq, test_label), verbose =
0.001 - val_loss: 0.0010 - val_mean_absolute_error: 0.0240
Epoch 7/100
28/28 [=====] - 1s 34ms/step - loss: 1.0998e-04 - mean_absolute_error:
0.0076 - val_loss: 0.0035 - val_mean_absolute_error: 0.0489
Epoch 8/100
28/28 [=====] - 1s 35ms/step - loss: 9.9768e-05 - mean_absolute_error:
0.0071 - val_loss: 0.0026 - val_mean_absolute_error: 0.0406
Epoch 9/100
28/28 [=====] - 1s 35ms/step - loss: 8.9201e-05 - mean_absolute_error:
0.0066 - val_loss: 0.0016 - val_mean_absolute_error: 0.0305
Epoch 10/100
28/28 [=====] - 1s 39ms/step - loss: 9.1312e-05 - mean_absolute_error:
0.0067 - val_loss: 0.0028 - val_mean_absolute_error: 0.0422
Epoch 11/100
28/28 [=====] - 1s 40ms/step - loss: 9.6708e-05 - mean_absolute_error:
0.0070 - val_loss: 0.0020 - val_mean_absolute_error: 0.0347
Epoch 12/100
28/28 [=====] - 1s 39ms/step - loss: 9.7572e-05 - mean_absolute_error:
0.0069 - val_loss: 0.0025 - val_mean_absolute_error: 0.0400
Epoch 13/100
```

```
In [21]: # predicting the values after running the model
test_predicted = model.predict(test_seq)
test_predicted[:5]

9/9 [=====] - 1s 12ms/step
```

```
Out[21]: array([[0.4098233, 0.4075212],
 [0.41005713, 0.40773788],
 [0.4094674, 0.407063],
 [0.41640925, 0.41348976],
 [0.41758883, 0.41485825]], dtype=float32)
```

```
In [22]: # Inversing normalization/scaling on predicted data
test_inverse_predicted = mms.inverse_transform(test_predicted)
test_inverse_predicted[:5]
```

```
Out[22]: array([[1430.7797, 1423.5354],
 [1431.2131, 1423.9369],
 [1430.1198, 1422.6862],
 [1442.9895, 1434.5972],
 [1445.1763, 1437.1334]], dtype=float32)
```

PREDICTED DATA VS VISUALIZING ACTUAL


```
In [23]: # Merging actual and predicted data for better visualization
df_merge = pd.concat([df.iloc[-264:].copy(),
                      pd.DataFrame(test_inverse_predicted, columns=['open_predicted', 'close_pred
                      index=df.iloc[-264:].index)], axis=1)
```

```
In [24]: # Inversing normalization/scaling
df_merge[['open', 'close']] = mms.inverse_transform(df_merge[['open', 'close']])
df_merge.head()
```


```
Out[24]:
```

	open	close	open_predicted	close_predicted
date				
2020-05-27	1417.25	1417.84	1430.779663	1423.535400
2020-05-28	1396.86	1416.73	1431.213135	1423.936890
2020-05-29	1416.94	1428.92	1430.119751	1422.686157
2020-06-01	1418.39	1431.82	1442.989502	1434.597168
2020-06-02	1430.55	1439.22	1445.176270	1437.133423

```
In [25]: # plotting the actual open and predicted open prices on date index
df_merge[['open', 'open_predicted']].plot(figsize=(10,6))
plt.xticks(rotation=45)
plt.xlabel('Date',size=15)
plt.ylabel('Stock Price',size=15)
plt.title('Actual vs Predicted for open price',size=15)
plt.show()
```



```
In [26]: # plotting the actual close and predicted close prices on date index
df_merge[['close', 'close_predicted']].plot(figsize=(10,6))
plt.xticks(rotation=45)
plt.xlabel('Date',size=15)
plt.ylabel('Stock Price',size=15)
plt.title('Actual vs Predicted for close price',size=15)
plt.show()
```



PREDICTING UPCOMING 10 DAYS

```
In [27]: # Creating a dataframe and adding 10 days to existing index

df_merge = df_merge.append(pd.DataFrame(columns=df_merge.columns,
                                         index=pd.date_range(start=df_merge.index[-1], periods=11, f
df_merge['2021-06-09':'2021-06-16']
```

```
Out[27]:
```

	open	close	open_predicted	close_predicted
2021-06-09	2499.50	2491.40	2334.801514	2328.246826
2021-06-10	2494.01	2521.60	2323.839600	2319.044189
2021-06-11	2524.92	2513.93	2364.799561	2357.900635
2021-06-12	NaN	NaN	NaN	NaN
2021-06-13	NaN	NaN	NaN	NaN
2021-06-14	NaN	NaN	NaN	NaN
2021-06-15	NaN	NaN	NaN	NaN
2021-06-16	NaN	NaN	NaN	NaN

```
In [28]: # creating a DataFrame and filling values of open and close column
upcoming_prediction = pd.DataFrame(columns=['open','close'],index=df_merge.index)
upcoming_prediction.index=pd.to_datetime(upcoming_prediction.index)
```


```
In [29]: curr_seq = test_seq[-1:]

for i in range(-10,0):
    up_pred = model.predict(curr_seq)
    upcoming_prediction.iloc[i] = up_pred
    curr_seq = np.append(curr_seq[0][1:],up_pred,axis=0)
    curr_seq = curr_seq.reshape(test_seq[-1:].shape)


1/1 [=====] - 0s 96ms/step
1/1 [=====] - 0s 209ms/step
1/1 [=====] - 0s 82ms/step
1/1 [=====] - 0s 51ms/step
1/1 [=====] - 0s 205ms/step
1/1 [=====] - 0s 356ms/step
1/1 [=====] - 0s 100ms/step
1/1 [=====] - 0s 169ms/step
1/1 [=====] - 0s 61ms/step
1/1 [=====] - 0s 41ms/step
```

```
In [30]: # inversing Normalization/scaling
upcoming_prediction[['open','close']] = mms.inverse_transform(upcoming_prediction[['open','close']])
```

```
In [31]: # plotting Upcoming Open price on date index
fig,ax=plt.subplots(figsize=(10,5))
ax.plot(df_merge.loc['2021-04-01:', 'open'],label='Current Open Price')
ax.plot(upcoming_prediction.loc['2021-04-01:', 'open'],label='Upcoming Open Price')
plt.setp(ax.xaxis.get_majorticklabels(), rotation=45)
ax.set_xlabel('Date',size=15)
ax.set_ylabel('Stock Price',size=15)
ax.set_title('Upcoming Open price prediction',size=15)
ax.legend()
fig.show()
```




```
In [32]: # plotting Upcoming Close price on date index
fig,ax=plt.subplots(figsize=(10,5))
ax.plot(df_merge.loc['2021-04-01':,'close'],label='Current close Price')
ax.plot(upcoming_prediction.loc['2021-04-01':,'close'],label='Upcoming close Price')
plt.setp(ax.xaxis.get_majorticklabels(), rotation=45)
ax.set_xlabel('Date',size=15)
ax.set_ylabel('Stock Price',size=15)
ax.set_title('Upcoming close price prediction',size=15)
ax.legend()
fig.show()
```



THANK YOU!

GitHub Link: <https://github.com/anujtiwari21?tab=repositories>
[\(https://github.com/anujtiwari21?tab=repositories\)](https://github.com/anujtiwari21?tab=repositories)